

Estimation Practicum

CAPM Replication

General Regression Techniques


- Fixed Effects
- Dealing with outliers via $\ln()$
- Clustering of standard errors
- Interaction and squared terms
- Collinearity
- Discrete dependent variables
- Probit, Logit, Multinomial and Ordered

Dataset and code same as for class 3. If you have not installed Python yet, make this a priority

CAPM Tests

Undervalued versus Overvalued

- From (5), we see that the market portfolio has $\beta_m = 1$.
- You can easily draw an SML with two points $(0, r_f)$ and $(1, r_m)$.





Implications and Tests of CAPM

- If **CAPM** were correct:
 - OLS regression of excess asset return against excess market returns for any asset **would have a constant term of** _____
 - OLS regression of excess asset return against excess market returns + **{any other variable on the RHS}** would have a {significant | insignificant} coefficient on the other variables
 - Above **two points help us to test CAPM empirically**
- Some **CAPM tests**:
 - Fama-French showed that in addition to market beta, company size (proxied by market cap) and valuation (book to market ratio) are also significant. This became **Fama-French “3 factor model”**
 - Carhart, proposed a 4th factor, **which is momentum, proxied by lagged returns of the stock**
 - Thereafter, the ‘flood gates’ opened, and **researchers started adding and testing additions using OLS framework**

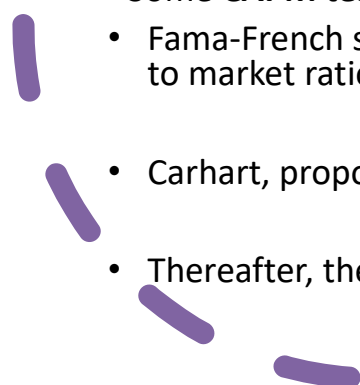


Table 1
Factor classification

	Risk type	Description	Examples
Common (113)	Financial (46)	Proxy for aggregate financial market movement, including market portfolio returns, volatility, squared market returns, among others	Sharpe (1964): market returns; Kraus and Litzenberger (1976): squared market returns
	Macro (40)	Proxy for movement in macroeconomic fundamentals, including consumption, investment, inflation, among others	Breeden (1979): consumption growth; Cochrane (1991): investment returns
	Microstructure (11)	Proxy for aggregate movements in market microstructure or financial market frictions, including liquidity, transaction costs, among others	Pastor and Stambaugh (2003): market liquidity; Lo and Wang (2006): market trading volume
	Behavioral (3)	Proxy for aggregate movements in investor behavior, sentiment or behavior-driven systematic mispricing	Baker and Wurgler (2006): investor sentiment; Hirshleifer and Jiang (2010): market mispricing
	Accounting (8)	Proxy for aggregate movement in firm-level accounting variables, including payout yield, cash flow, among others	Fama and French (1992): size and book-to-market; Da and Warachka (2009): cash flow
	Other (5)	Proxy for aggregate movements that do not fall into the above categories, including momentum, investors' beliefs, among others	Carhart (1997): return momentum; Ozoguz (2009): investors' beliefs
Characteristics (202)	Financial (61)	Proxy for firm-level idiosyncratic financial risks, including volatility, extreme returns, among others	Ang et al. (2006): idiosyncratic volatility; Bali, Cakici, and Whitelaw (2011): extreme stock returns
	Microstructure (28)	Proxy for firm-level financial market frictions, including short sale restrictions, transaction costs, among others	Jarrow (1980): short sale restrictions; Mayshar (1981): transaction costs
	Behavioral (3)	Proxy for firm-level behavioral biases, including analyst dispersion, media coverage, among others	Diether, Malloy, and Scherbina (2002): analyst dispersion; Fang and Peress (2009): media coverage
	Accounting (87)	Proxy for firm-level accounting variables, including PE ratio, debt-to-equity ratio, among others	Basu (1977): PE ratio; Bhandari (1988): debt-to-equity ratio
	Other (24)	Proxy for firm-level variables that do not fall into the above categories, including political campaign contributions, ranking-related firm intangibles, among others	Cooper, Gulen, and Ovtchinnikov (2010): political campaign contributions; Edmans (2011): intangibles

The numbers in parentheses represent the number of factors identified. See Table 6 and <http://faculty.fuqua.duke.edu/~charvey/Factor-List.xlsx>.

Currently, there are **around 300 factors** published in top tier academic journals:
<https://academic.oup.com/rfs/article/29/1/5/1843824>

Application of CAPM: Cross Sectional Stock Picking

- Ideal stock to hold:
 - positive alpha
 - large beta during bull market
 - small beta during bear market
 - large Sharpe ratio
- Ideal stock to short is the reverse.
- Long the “good” stocks, short the “bad” stocks. Will this quant strategy work?
- So exactly how could one search for those good and bad stocks?


Compare Cross Sectional Stock Picking with Market Timing

- Shift investment funds into the market portfolio when market is rising, and to shift out of the stock market into money market when market is falling, especially if the market falls below risk-free return.
- The goal is to avoid being invested in e.g. mutual funds during a market decline.
- Typically, trend-following indicators are used to determine the direction and identify buy and sell signals.
- In an up move “buy signal,” money is transferred from a money market fund into a mutual fund in an attempt to capture a capital gain.
- In a down move “sell signal,” the assets in the mutual fund are sold and moved back into the money market for safe keeping until the next up move.

Jensen's Empirical Tests


- Testing for the presence and significance of abnormal returns (“Jensen’s alpha” - Jensen, 1968).
- Data: Annual Returns on the portfolios of 115 mutual funds from 1945–1964.
- The model: $R_{jt} - R_{ft} = \alpha_j + \beta_j (R_{mt} - R_{ft}) + u_{jt}$ for $j = 1, 2, \dots, 115$.
- Are α_j significant?
- The null hypothesis is $H_0: \alpha_j = 0$.

Distribution of Mutual Fund Alphas' t-Ratios




Jensen (1968). Reprinted with the permission of Blackwell publishers.

Net of Transactions Costs



Jensen (1968). Reprinted with the permission of Blackwell publishers.

General Regression Techniques



In this section, we implement some regressions in python using both OLS as well as logit models, as well as explore some practical issues with outliers, fixed effects, and standard error estimation



Python packages used are fully open source (pandas, numpy, statsmodels), and are also widely used in industry or research academia.



There are many python versions. One convenient version that will allow you to follow along is here (<https://www.anaconda.com/products/individual>). Notwithstanding, we make no judgement that any version of python is better or worse; if you already have it installed, just use pre-existing



Datasets and sample code

Same as for Class 3

- On E-Learn:
 - **corpfund.csv** contains panel data on corporate performance (e.g. EBITDA, net income, revenue, etc) for US listed firms
 - **stockmetadata.csv** contains company descriptions such as tickers, industry, geographical location, etc
 - **ols_tests.py** contains python code that ingests and merges the data, as well as runs statistical analysis
 - Do not share data with anyone who is not taking this class

Data ingestion and merging steps

```
1 import pandas as pd
2 import numpy as np
3 import statsmodels.api as sm
4 import statsmodels.discrete.discrete_model as smdiscrete
5 pd.set_option('use_inf_as_na', True)
6
7 meta_df = pd.read_csv("stockmetadata.csv")
8 fdata_df = pd.read_csv("corpfund.csv")
9 fdata_df = fdata_df[fdata_df['dimension'] == 'ARQ']
10 fdata_df['datekey'] = pd.to_datetime(fdata_df['datekey'])
11 df_left = pd.merge(fdata_df, meta_df, on='ticker', how='left')
12 df_left = df_left.set_index('datekey')
```

1. Lines 7 & 8: read in both csv files as python dataframes
2. A python dataframe is a table with (usually) time on the x-axis and various variables in columns
3. In line 9, we only keep “as reported data” to avoid forward bias due to corporate restatements
4. In line 11, we merge both dataframes based on ticker
5. In line 12, we set index of the dataframe to be data-date

```

In [2]: df_left
Out[2]:

```

datekey	ticker	dimension	calendardate	reportperiod	lastupdated_x	accoci	...	firstpricedate	lastpricedate	firstquarter	lastquarter	secfilings	companysite
2010-06-07	A	ARQ	2010-03-31	2010-04-30	2020-09-01	-239000000.0	...	1999-11-18	2020-10-09	1997-06-30	2020-06-30	https://www.sec.gov/cgi-bin/browse-edgar?actio...	http://www.agilent.com
2010-09-07	A	ARQ	2010-06-30	2010-07-31	2020-09-01	-225000000.0	...	1999-11-18	2020-10-09	1997-06-30	2020-06-30	https://www.sec.gov/cgi-bin/browse-edgar?actio...	http://www.agilent.com
2010-12-20	A	ARQ	2010-09-30	2010-10-31	2020-09-01	-880000000.0	...	1999-11-18	2020-10-09	1997-06-30	2020-06-30	https://www.sec.gov/cgi-bin/browse-edgar?actio...	http://www.agilent.com
2011-03-09	A	ARQ	2010-12-31	2011-01-31	2020-09-01	-630000000.0	...	1999-11-18	2020-10-09	1997-06-30	2020-06-30	https://www.sec.gov/cgi-bin/browse-edgar?actio...	http://www.agilent.com
2011-06-07	A	ARQ	2011-03-31	2011-04-30	2020-09-01	278000000.0	...	1999-11-18	2020-10-09	1997-06-30	2020-06-30	https://www.sec.gov/cgi-bin/browse-edgar?actio...	http://www.agilent.com
...
2020-10-05	STTK	ARQ	2019-12-31	2019-12-31	2020-10-09	54000.0	...	2020-10-09	2020-10-09	2018-12-31	2020-06-30	https://www.sec.gov/cgi-bin/browse-edgar?actio...	http://www.shattucklabs.com
2020-10-08	STTK	ARQ	2020-06-30	2020-06-30	2020-10-09	18000.0	...	2020-10-09	2020-10-09	2018-12-31	2020-06-30	https://www.sec.gov/cgi-bin/browse-edgar?actio...	http://www.shattucklabs.com
2020-10-08	KRON	ARQ	2020-06-30	2020-06-30	2020-10-09	164000.0	...	2020-10-09	2020-10-09	2018-12-31	2020-06-30	https://www.sec.gov/cgi-bin/browse-edgar?actio...	http://kronosbio.com
2020-09-18	SPRB	ARQ	2019-12-31	2019-12-31	2020-10-09	0.0	...	2020-10-09	2020-10-09	2018-12-31	2020-06-30	https://www.sec.gov/cgi-bin/browse-edgar?actio...	http://www.sprucebiosciences.com
2020-10-05	SPRB	ARQ	2020-06-30	2020-06-30	2020-10-09	0.0	...	2020-10-09	2020-10-09	2018-12-31	2020-06-30	https://www.sec.gov/cgi-bin/browse-edgar?actio...	http://www.sprucebiosciences.com

```

[207494 rows x 138 columns]

```

This is what dataset looks like after running code from previous slide

Fixed effects

- We want to estimate relationship between E/P ratio and operating profit margin for each of these data point
- However, we note different industries may have completely different operating margins on average, as well as different relationships between operating margin and E/P ratio
- One solution is to introduce an indicator variable (“dummy variable”) for each industry
- This allows each industry to effectively have its own y-intercept

Fixed effects specification

OLS without fixed effects:

$$Y_i = \alpha + \beta X_i + \varepsilon$$

OLS with fixed effects:

$$Y_i = \alpha + \beta_1 X_i + \beta_2 * I_1 + \beta_3 * I_2 + \beta_4 * I_3 + \dots \beta_{12} * I_{12} + \varepsilon$$

Where

$I_1 = 1$ if firm is in sector 1, 0 otherwise

$I_2 = 1$ if firm is in sector 2, 0 otherwise

...

$I_{12} = 1$ if firm is in sector 12, 0 otherwise

Consequently:

Effective y-intercept for datapoints in sector 1 = $\alpha + \beta_2$

Effective y-intercept for datapoints in sector 2 = $\alpha + \beta_3$, etc

```

15 industrydummies = pd.get_dummies(df_left['sicsector'])
16 industrydummies.sum()           #purely for exploring the data, has no other purpose
17 industrydummies.describe()      #purely for exploring the data, has no other purpose
18
19 data_w_dummies = pd.concat([df_left, industrydummies], axis=1)
20
21 data_w_dummies.drop(['Wholesale Trade'], inplace=True, axis=1) #drop 1 dummy variable
22 data_w_dummies['epratio'] = data_w_dummies['eps']/data_w_dummies['price'] #generate dependent variable
23 data_w_dummies['operatingmargin'] = data_w_dummies['opinc'] / data_w_dummies['revenue'] #generate independent variable

```

Fixed effects example code

- Line 15, we generate 1 dummy variable for each sector
- Line 19, we merge the dummy variable dataframe into the main dataframe
- Line 21, we drop 1 dummy variable, which will be the control group
- Line 22, we generate our dependent variable (E/P ratio)
- Line 23, we generate our independent variable, which is operating margin

```

In [3]: industrydummies
Out[3]:
datekey  Agriculture Forestry And Fishing  Construction  Finance Insurance And Real Estate  Manufacturing  Mining  Retail Trade  Services  Transportation Communications Electric Gas And Sanitary Service  Wholesale Trade
2010-06-07                                0              0                                0              1      0              0      0              0              0
2010-09-07                                0              0                                0              1      0              0      0              0              0
2010-12-20                                0              0                                0              1      0              0      0              0              0
2011-03-09                                0              0                                0              1      0              0      0              0              0
2011-06-07                                0              0                                0              1      0              0      0              0              0
...
2020-10-05                                0              0                                0              1      0              0      0              0              0
2020-10-08                                0              0                                0              1      0              0      0              0              0
2020-10-08                                0              0                                0              1      0              0      0              0              0
2020-09-18                                0              0                                0              1      0              0      0              0              0
2020-10-05                                0              0                                0              1      0              0      0              0              0
[207494 rows x 9 columns]

```

This is what ‘industrydummies’ dataframe looks like after line 15 in the previous slide

```

26 #initial analysis
27 result = sm.OLS(data_w_dummies['epratio'], sm.add_constant(data_w_dummies[['operatingmargin']], missing='drop')).fit()
28 result.summary()
29 result = sm.OLS(data_w_dummies['epratio'], sm.add_constant(data_w_dummies[['operatingmargin', 'Agriculture Forestry And Fishing', 'Construction', 'Finance Insurance And Real Estate',
30 result.summary()

```

OLS regression with and without dummy variables

- Line 27 and 28, we run the OLS regression without dummy variables
- Line 29 and 30, we include the dummy variables

```

In [4]: result = sm.OLS(data_w_dummies['epratio'], sm.add_constant(data_w_dummies[['operatingmargin']])),
...: result.summary()
Out[4]:
<class 'statsmodels.iolib.summary.Summary'>
"""
                        OLS Regression Results
=====
Dep. Variable:          epratio      R-squared:                0.000
Model:                  OLS          Adj. R-squared:           -0.000
Method:                 Least Squares   F-statistic:             4.803e-06
Date:                   Sat, 10 Oct 2020   Prob (F-statistic):      0.998
Time:                   16:58:52          Log-Likelihood:          -6.4512e+05
No. Observations:       115861           AIC:                    1.290e+06
Df Residuals:           115859           BIC:                    1.290e+06
Df Model:                1
Covariance Type:        nonrobust
=====
                        coef      std err          t      P>|t|      [0.025      0.975]
-----
const                0.2554      0.186        1.372      0.170      -0.110      0.620
operatingmargin    4.402e-06      0.002        0.002      0.998      -0.004      0.004
=====
Omnibus:                 651923.203    Durbin-Watson:           1.999
Prob(Omnibus):            0.000    Jarque-Bera (JB):    60232700557916.930
Skew:                    331.783    Prob(JB):              0.00
Kurtosis:                111701.024    Cond. No.              92.7
=====

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

No Fixed Effects

```

...: result.summary()
Out[5]:
<class 'statsmodels.iolib.summary.Summary'>
"""
                        OLS Regression Results
=====
Dep. Variable:          epratio    R-squared:                0.000
Model:                  OLS      Adj. R-squared:             -0.000
Method:                 Least Squares    F-statistic:          0.3172
Date:                   Sat, 10 Oct 2020    Prob (F-statistic):      0.970
Time:                   16:59:55    Log-Likelihood:         -6.4512e+05
No. Observations:      115861    AIC:                   1.290e+06
Df Residuals:          115851    BIC:                   1.290e+06
Df Model:               9
Covariance Type:       nonrobust
=====
                        coef    std err          t      P>|t|      [0.025     0.975]
-----
const                0.7960      0.372        2.141      0.032      0.067      1.525
operatingmargin      -1.244e-05      0.002       -0.006      0.995     -0.004      0.004
Agriculture Forestry And Fishing
Construction         -0.7124      3.933       -0.181      0.856     -8.422      6.997
Finance Insurance And Real Estate
Manufacturing        -0.7724      1.732       -0.446      0.656     -4.168      2.623
Mining               -0.6747      0.531       -1.270      0.204     -1.716      0.366
Retail Trade         -0.7171      0.525       -1.367      0.172     -1.745      0.311
Services             -0.7509      1.226       -0.612      0.540     -3.154      1.653
Transportation Communications Electric Gas And Sanitary Service
Retail Trade         -0.7757      0.927       -0.837      0.402     -2.592      1.040
Services            -0.7737      0.689       -1.122      0.262     -2.125      0.577
Transportation Communications Electric Gas And Sanitary Service
=====
Omnibus:              651916.510    Durbin-Watson:          1.999
Prob(Omnibus):        0.000    Jarque-Bera (JB):      60226461523137.062
Skew:                 331.771    Prob(JB):              0.00
Kurtosis:             111695.238    Cond. No.              1.96e+03
=====

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 1.96e+03. This might indicate that there are
strong multicollinearity or other numerical problems.
"""

```

With Fixed Effects

We note t-statistic for estimated coefficient on 'operatingmargin' is very low

This might be due to the effect of outliers overly influencing estimation

Two ways to explore the effect of outliers are:

- Winsorize data [albeit at expense of deleting observations arbitrarily ("datamining")]
- Apply logarithm to both dependent and independent variables. This will naturally reduce impact of outliers

```

31 data_w_dummies['lnoperatingmargin'] = np.log(data_w_dummies['operatingmargin'])
32 result = sm.OLS(data_w_dummies['epratio'], sm.add_constant(data_w_dummies[['lnoperatingmargin']], missing='drop')).fit()
33 result.summary()
34 data_w_dummies['lnepratio'] = np.log(data_w_dummies['epratio'])
35 result = sm.OLS(data_w_dummies['lnepratio'], sm.add_constant(data_w_dummies[['lnoperatingmargin']], missing='drop')).fit()
36 result.summary()
37
38 #with dummy variables
39 result = sm.OLS(data_w_dummies['lnepratio'], sm.add_constant(data_w_dummies[['lnoperatingmargin', 'Agriculture Forestry And Fishing', 'Construction', 'Finance Insurance And Real Estate']])).fit()
40 result.summary()

```

- Lines 31 and 34: Generate log versions of both dependent and independent variables
- Line 35: Run regression with log variables
- Line 39: Same regression, but with dummy variables included


```

In [6]: result = sm.OLS(data_w_dummies['lnepratio'], sm.add_constant(data_w_dummies[['lnoperatingmargin']]), missing='drop').fit()
...: result.summary()^M
...:
Out[6]:
<class 'statsmodels.iolib.summary.Summary'>
"""
                        OLS Regression Results
=====
Dep. Variable:          lnepratio    R-squared:                0.071
Model:                OLS          Adj. R-squared:            0.071
Method:             Least Squares   F-statistic:             8846.
Date:               Sat, 10 Oct 2020 Prob (F-statistic):        0.00
Time:               17:08:40         Log-Likelihood:         -1.5681e+05
No. Observations:    115861         AIC:                   3.136e+05
Df Residuals:        115859         BIC:                   3.136e+05
Df Model:             1
Covariance Type:     nonrobust
=====
                        coef    std err          t      P>|t|      [0.025      0.975]
-----
const                -3.7721      0.006   -602.657      0.000     -3.784     -3.760
lnoperatingmargin      0.2638      0.003    94.054      0.000      0.258      0.269
=====
Omnibus:            18365.631   Durbin-Watson:           1.036
Prob(Omnibus):        0.000   Jarque-Bera (JB):       158342.851
Skew:                0.510   Prob(JB):                0.00
Kurtosis:            8.635   Cond. No.                5.93
=====

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
"""

```

Log variables, no FE

```

In [7]: result = sm.OLS(data_w_dummies['lnepratio'], sm.add_constant(data_w_dummies[['lnepratio', 'Agriculture Forestry And Fishi
...: 'Services', 'Transportation Communications Electric Gas And Sanitary Service']]), missing='drop').fit()
...: result.summary()
Out[7]:
<class 'statsmodels.iolib.summary.Summary'>
"""
                        OLS Regression Results
=====
Dep. Variable:          lnepratio      R-squared:                0.085
Model:                  OLS           Adj. R-squared:            0.084
Method:                 Least Squares   F-statistic:              1188.
Date:                   Sat, 10 Oct 2020   Prob (F-statistic):       0.00
Time:                   17:09:31          Log-Likelihood:           -1.5595e+05
No. Observations:       115861          AIC:                     3.119e+05
Df Residuals:           115851          BIC:                     3.120e+05
Df Model:                9
Covariance Type:        nonrobust
=====
                        coef      std err          t      P>|t|      [0.025      0.975]
-----
const                -3.6332      0.009    -419.187      0.000     -3.650     -3.616
lnepratio              0.2969      0.003     96.466      0.000      0.291      0.303
Agriculture Forestry And Fishing
0.6587      0.058     11.417      0.000      0.546      0.772
Construction          0.3404      0.025     13.352      0.000      0.290      0.390
Finance Insurance And Real Estate
-0.1784      0.008    -21.688      0.000     -0.195     -0.162
Manufacturing         -0.0182      0.008     -2.360      0.018     -0.033     -0.003
Mining                -0.0450      0.018     -2.489      0.013     -0.080     -0.010
Retail Trade           0.1256      0.014      9.187      0.000      0.099      0.152
Services              -0.2630      0.010    -26.005      0.000     -0.283     -0.243
Transportation Communications Electric Gas And Sanitary Service
-0.1285      0.012    -10.640      0.000     -0.152     -0.105
=====
Omnibus:             18757.262   Durbin-Watson:           1.031
Prob(Omnibus):        0.000   Jarque-Bera (JB):        159948.347
Skew:                 0.533   Prob(JB):                0.00
Kurtosis:             8.657   Cond. No.                51.5
=====

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
"""

```

Log variables, with FE

Clustering of Standard Errors



We note t-statistic is significantly improved, but could this be a result of intra-cluster correlations?

- Observations within same industry are likely to be highly correlated versus observations in different industries
- This violates OLS assumption of iid datapoints
- Estimates may therefore greatly overstate efficiency (inflated t-statistics, artificially low p-values, downward biases standard errors)
- We can adjust the standard error estimates by clustering errors within each industry
- Background: See for e.g. <https://economics.mit.edu/files/13927> for more information

For next couple of slides, we focus on empirical application rather than theoretical proof of forgoing

```

42 #clustering standard variables
43 data_w_dummies.dropna(subset=['lnepratio', 'lnoperatingmargin'], inplace=True) #because of a bug in python where fillna is not working perfectly
44 #note that we can cannot cluster by str variables in python, hence using siccode instead of sicsector
45 result = sm.OLS(data_w_dummies['lnepratio'], sm.add_constant(data_w_dummies[['lnoperatingmargin', 'Agriculture Forestry And Fishing', 'Construction', 'Finance Insurance And Real Es
&Trade', 'Services', 'Transportation Communications Electric Gas And Sanitary Service']]), missing='drop').fit(cov_type='cluster', cov_kwds={'groups': data_w_dummies['siccode']})
46 result.summary()
47

```



Line 43: We drop all observations with missing data to avoid a bug in python



Line 45: Run OLS estimation on log variables with FE, and standard errors clustered by sic_code

```

In [8]: result = sm.OLS(data_w_dummies['lnepratio'], sm.add_constant(data_w_dummies[['lnoperatingmargin', 'Agriculture Forestry And Fishing',
...: 'Services', 'Transportation Communications Electric Gas And Sanitary Service']]), missing='drop').fit(cov_type='cluster', cov_kwds=
...: result.summary()
...:
Out[8]:
<class 'statsmodels.iolib.summary.Summary'>
"""
                        OLS Regression Results
=====
Dep. Variable:          lnepratio    R-squared:                0.085
Model:                  OLS          Adj. R-squared:            0.084
Method:                 Least Squares   F-statistic:              44.65
Date:                   Sat, 10 Oct 2020   Prob (F-statistic):       2.30e-80
Time:                   17:22:14          Log-Likelihood:           -1.5595e+05
No. Observations:       115861          AIC:                     3.119e+05
Df Residuals:           115851          BIC:                     3.120e+05
Df Model:                9
Covariance Type:        cluster
=====
                        coef      std err          z      P>|z|      [0.025      0.975]
-----
const                  -3.6332      0.034    -107.420      0.000      -3.700      -3.567
lnoperatingmargin       0.2969      0.016     18.678      0.000       0.266       0.328
Agriculture Forestry And Fishing
Construction            0.6587      0.357      1.847      0.065      -0.040       1.358
Finance Insurance And Real Estate
Manufacturing           0.3404      0.064      5.292      0.000       0.214       0.467
Mining                  -0.1784      0.127     -1.405      0.160      -0.427       0.070
Retail Trade            -0.0182      0.052     -0.351      0.726      -0.120       0.083
Services                 0.1256      0.074      1.699      0.089      -0.019       0.270
Transportation Communications Electric Gas And Sanitary Service
                        -0.2630      0.085     -3.105      0.002      -0.429      -0.097
Transportation Communications Electric Gas And Sanitary Service
                        -0.1285      0.058     -2.215      0.027      -0.242      -0.015
=====
Omnibus:               18757.262    Durbin-Watson:           1.031
Prob(Omnibus):         0.000    Jarque-Bera (JB):        159948.347
Skew:                  0.533    Prob(JB):                0.00
Kurtosis:              8.657    Cond. No.                51.5
=====

Warnings:
[1] Standard Errors are robust to cluster correlation (cluster)

```

Estimated t-statistic is several times lower after clustering!

Interaction Variables and Squared Effects

Interaction variables

- The effect of an independent variable may depend on the value of another independent variable in the system
- For instance, a variable on profitability may have a different predictive effect on stock price depending on whether the company's revenue is growing or not
- Interaction variables allow us to model different slope coefficients in linear systems for different situations
- It is similar to dummy variables, where you can model different intercepts for various situations


```

99 #####INTERACTION_VARIABLES###
100 dummies = industrydummies.columns.values
101 dummies = dummies[1:len(dummies)-1]
102 dlist = []
103 for ind in dummies:
104     tag = 'lnoperatingmargin'+ '_' +ind
105     data[tag] = data['lnoperatingmargin']*(data[ind])
106     dlist.append(tag)
107 dlist.append(ind)
108 dlist.append('lnoperatingmargin')
109 result = sm.OLS(data['lnopratio'], sm.add_constant(data[dlist]), missing='drop').fit(cov_type='cluster', cov_kwds={'groups': data_w_dummies['siccode']})
110 result.summary()

```

Lines 99 to 107: Generate the interaction variables. In this case, we create interaction variables as lnoperatingmargin*industry dummy

Lines 108 to 110: Run OLS estimation with interaction variables

```

=====
                        OLS Regression Results
=====
Dep. Variable:          lnopratio    R-squared:                0.053
Model:                  OLS         Adj. R-squared:           0.052
Method:                 Least Squares   F-statistic:             22.55
Date:                   Fri, 16 Oct 2020   Prob (F-statistic):      2.45e-61
Time:                   20:01:30         Log-Likelihood:          -41725.
No. Observations:       31998          AIC:                    8.348e+04
Df Residuals:           31982          BIC:                    8.361e+04
Df Model:                15
Covariance Type:        cluster
=====

```

	coef	std err	z	P> z	[0.025	0.975]
const	-3.6108	0.048	-75.842	0.000	-3.704	-3.518
lnoperatingmargin_Construction	0.2473	0.109	2.266	0.023	0.033	0.461
Construction	0.7646	0.378	2.021	0.043	0.023	1.506
lnoperatingmargin_Finance Insurance And Real Estate	-0.1840	0.194	-0.947	0.343	-0.565	0.197
Finance Insurance And Real Estate	-0.8341	0.417	-2.002	0.045	-1.651	-0.018
lnoperatingmargin_Manufacturing	-0.0364	0.066	-0.555	0.579	-0.165	0.092
Manufacturing	-0.2274	0.106	-2.151	0.031	-0.435	-0.020
lnoperatingmargin_Mining	0.3526	0.084	4.213	0.000	0.189	0.517
Mining	0.4767	0.143	3.334	0.001	0.196	0.757
lnoperatingmargin_Retail Trade	-0.0927	0.082	-1.132	0.258	-0.253	0.068
Retail Trade	-0.2019	0.218	-0.926	0.354	-0.629	0.225
lnoperatingmargin_Services	0.1619	0.060	2.720	0.007	0.045	0.279
Services	-0.0084	0.132	-0.063	0.950	-0.268	0.251
lnoperatingmargin_Transportation Communications Electric Gas And Sanitary Service	0.0424	0.075	0.567	0.571	-0.104	0.189
Transportation Communications Electric Gas And Sanitary Service	-0.0560	0.163	-0.343	0.732	-0.376	0.264
lnoperatingmargin	0.2399	0.021	11.437	0.000	0.199	0.281

```

=====
Omnibus:                8577.755    Durbin-Watson:           1.016
Prob(Omnibus):          0.000      Jarque-Bera (JB):        89878.807
Skew:                   0.988      Prob(JB):                0.00
Kurtosis:               10.969     Cond. No.:               144.
=====
Warnings:
[1] Standard Errors are robust to cluster correlation (cluster)
=====

```

What does the coefficient of each interaction variable mean?

Approximating non-linear relationships with OLS estimations: Squared and higher order terms



We may have convex or concave relationships in finance / economics (e.g. utility functions, industrial cost functions ...)



A convex or concave relationship between dependent and independent variable can be modelled by adding squared independent variables to the estimation



Higher order terms are also possible if this is motivated by theory

Note that using log independent variables also models non-linear relationships

```

114 #####SQUARED VARIABLES#####
115 #We hypothesize that a stock's valuation has a concave relationship to its asset tangibility - why?
116 data['tangibles_assets_normal'] = data['tangibles']/data['assets']
117 data['tangibles_assets_sq'] = data['tangibles_assets_normal']*data['tangibles_assets_normal']
118 result = sm.OLS(data['epratio'], sm.add_constant(data[['tangibles_assets_normal', 'tangibles_assets_sq']])), missing='drop').fit(cov_type='cluster', cov_kws={'groups': da
119 print(result.summary())
120 result = sm.OLS(data['epratio'], sm.add_constant(data[['tangibles_assets_normal']])), missing='drop').fit(cov_type='cluster', cov_kws={'groups': data_w_dummies['siccode']
121 print(result.summary())

```

- We explore in lines 114 to 121 the hypothesis that a stock's earnings may have a concave relationship with its asset tangibility (plant property equipment / total assets)
- Initially, having more asset tangibility could allow the firm greater access to financing, because the assets can be pledged as collateral
- However, there is an optimal point ("maxima"). Beyond that, having too much fixed assets may mean the firm has less cash on its balance sheet, giving it less freedom to pursue growth opportunities or react to emergencies
- Lines 116 to 117: Generate measures of asset tangibility and squared asset tangibility
- Lines 118 and 120: Run OLS versus only asset tangibility, versus against asset tangibility and asset tangibility_sq

```

Warnings:
[1] Standard Errors are robust to cluster correlation (cluster)
      OLS Regression Results
=====
Dep. Variable:          epratio      R-squared:                0.000
Model:                  OLS          Adj. R-squared:           -0.000
Method:                 Least Squares  F-statistic:              0.04372
Date:                   Fri, 16 Oct 2020  Prob (F-statistic):      0.834
Time:                   21:42:21      Log-Likelihood:          -1.0468e+05
No. Observations:       31998        AIC:                    2.094e+05
Df Residuals:           31996        BIC:                    2.094e+05
Df Model:                1
Covariance Type:        cluster
=====
               coef      std err          z      P>|z|      [0.025      0.975]
-----
const                0.0881      0.123      0.715      0.474      -0.153      0.329
tangibles_assets_normal -0.0258      0.123     -0.209      0.834      -0.267      0.216
=====
Omnibus:              151537.582    Durbin-Watson:           1.998
Prob(Omnibus):         0.000    Jarque-Bera (JB):    1359216358122.066
Skew:                  178.596    Prob(JB):             0.00
Kurtosis:              31930.236    Cond. No.              7.73
=====

Warnings:
[1] Standard Errors are robust to cluster correlation (cluster)

```

A univariate regression masks a more non-linear relationship ...

```

=====
                        OLS Regression Results
=====
Dep. Variable:          epratio    R-squared:                0.000
Model:                  OLS        Adj. R-squared:            -0.000
Method:                 Least Squares    F-statistic:              1.743
Date:                   Fri, 16 Oct 2020    Prob (F-statistic):       0.175
Time:                   21:42:21          Log-Likelihood:           -1.0468e+05
No. Observations:       31998            AIC:                     2.094e+05
Df Residuals:           31995            BIC:                     2.094e+05
Df Model:                2
Covariance Type:        cluster
=====
                        coef      std err          z      P>|z|      [0.025      0.975]
-----
const                -0.1497      0.245      -0.612      0.540      -0.629      0.330
tangibles_assets_normal  0.7366      1.058       0.696      0.486      -1.338      2.811
tangibles_assets_sq    -0.5502      0.850      -0.647      0.517      -2.216      1.116
=====
Omnibus:              151535.712    Durbin-Watson:           1.998
Prob(Omnibus):         0.000    Jarque-Bera (JB):    1359061937976.077
Skew:                  178.589    Prob(JB):             0.00
Kurtosis:              31928.422    Cond. No.              57.2
=====

Warnings:
[1] Standard Errors are robust to cluster correlation (cluster)

```

Which we can uncover with a squared term!

Multicollinearity

Multicollinearity

- Multicollinearity occurs when the explanatory variables are very highly correlated with each other.
- Perfect multicollinearity
 - Cannot estimate all the coefficients
 - e.g. suppose $x_3 = 2x_2$
 - and the model is $y_t = \beta_1 + \beta_2 x_{2t} + \beta_3 x_{3t} + \beta_4 x_{4t} + u_t$
- Problems if near multicollinearity is present
 - R^2 will be high but the individual coefficients will have high standard errors.
 - The regression becomes very sensitive to small changes in the specification.
 - Thus confidence intervals for the parameters will be very wide, and significance tests might therefore give inappropriate conclusions.

Measuring Multicollinearity

- The easiest way to measure the extent of multicollinearity is simply to look at the matrix of correlations between the individual variables. e.g.

corr	x_2	x_3	x_4
x_2	—	0.2	<u>0.8</u>
x_3	0.2	—	0.3
x_4	<u>0.8</u>	0.3	—

- But another problem: if 3 or more variables are linear
– e.g. $x_{2t} + x_{3t} = x_{4t}$,
then this method will not work.
- Note that high correlation between y and one of the x 's is not multicollinearity.

Solutions to the Problem of Multicollinearity

Some econometricians argue that if the model is otherwise OK, just ignore it.

Other ways to “cure” the problems are

- drop one of the collinear variables
- transform the highly correlated variables into a ratio
- go out and collect more data e.g.
 - a longer run of data
 - switch to a higher frequency

Challenge or opportunity?

- In our example dataset, there are hundreds of variables.
- This presents two challenges:
 - First, collinearity.
 - Second, it is difficult to interpret the result of extremely high dimensional estimations
- However, the number of variables also allows us to test a larger number of possible hypotheses
- Principal component analysis and/or clustering the data may allow us to leverage the strengths of the dataset while avoiding the above two challenges
- We will detail both of these techniques in class 8 on machine learning and classification



3. Discrete dependent variables

OLS is not the only widely used estimation algorithm

Probit / logit models


Used for binary dependent variables

Multinomial probit / logit

Categorical dependent variables

Ordered probit / logit

Discrete dependent variables




Probit / logit models

$Y_i = \alpha + \beta X_i$ where Y is either 1 or 0

Estimation options:


1. Just run OLS; estimate will still be unbiased and BLUE
2. Use a probit or logit model, which exploits fact that Y_i is binary.

Fitting a linear line to a dichotomous variable



Estimated line can continue infinitely above or below $[0,1]$


In this case, a
nonlinear curve
may suit the
data better



What curve have we seen in this class which is:

1. Naturally between $[0,1]$
2. Takes S-shape

CDF of normal distribution is used to determine probability of observing a 1 in the probit model



Estimating probit/logit models

Probit models are estimated using maximum likelihood (MLE), which is an iterative computational process

- Tries different values of coefficients repeatedly to estimate the parameters of a probability distribution
- Objective function is to maximize the likelihood of observing the dataset
- For a probit model, a normal distribution is used for the CDF, while in a logit function is used in the case of the logit model
- Density function of the logit model is very similar to standard normal (it may be computationally more efficient), but with thinner tails

Interpreting output of a logit model

01

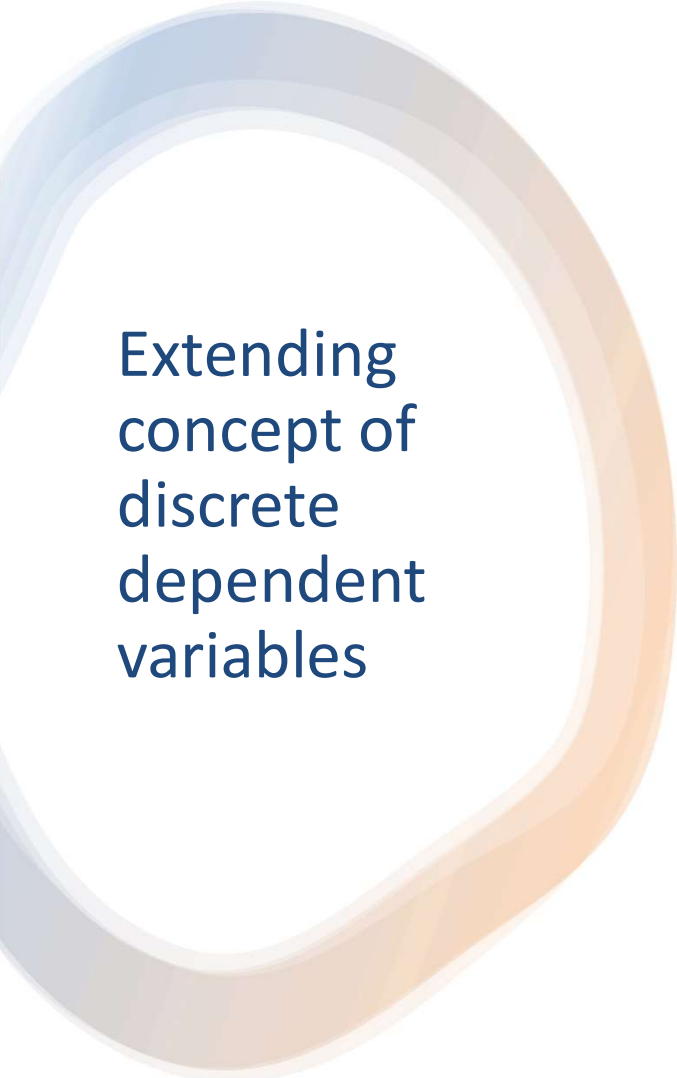
Estimated coefficients for a logit model are $\log(\text{odds-ratio})$

02

Odds-ratio of an event is $p(1-p)$ where p is the probability of the event

03

For the logit model, p is the probability of observing a 1



Extending concept of discrete dependent variables


- Returning to example of bond ratings, recall we can have 21 bond ratings e.g. with Moodys
- We can code each bond rating as “1”, “2”, “3”, etc in order of their probability of default. E.g. “1” corresponds to “AAA”, “21” to “CCC” etc
- Note there will be a natural ordering, where “1” is “better” than “2” in terms of default probability, etc



Ordered logit/probit models are a more general case of logit/probit

- Both ordered logit / probit models are a form of classification algorithm where the dependent value exists on an arbitrary scale, where only the relative ordering between different values is significant
- E.g. to model “customer satisfaction” from “5 – very good, 4 – good, 3 – neutral, etc”
- Estimation similar to logit/probit models, except multiple breakpoints on x-axis are estimated instead of just 1

Estimation of
ordered
logit/probit
model by
MLE



What if
dependent
variable is
purely
categorical?

- It is also possible for dependent variable to be purely discrete and categorical
- In this case, each value of dependent variable is merely a label, and larger or smaller values of dependent variable have no meaning
- For example, we may encode a discrete dependent variable to take values of 1, 2, 3, 4, 5, where 1 = blue, 2 = red, 3 = green, 4 = yellow and 5 = black
- In this case, $1 < 2 < 3 < \dots$ does not mean anything. The numbers are just **class labels**

We may use multinomial logit and probit models for classification problems

We wish to forecast which class / category a dependent variable will take, given independent variable values

Econometric methodology is to construct a score function that constructs a score from a set of weights that are linearly combined with the explanatory variables (features)

$$\text{score}(\mathbf{X}_i, k) = \beta_k \cdot \mathbf{X}_i$$

Multinomial estimation pseudo- algorithm

- Estimate coefficients for each score function independently (number of score functions is # categories - 1)
- Each score function gives (probability of an outcome) / (probability of base case)
- All probabilities must sum to 1, which gives us probability of base case
- Pick alternative with highest probability
- Estimation framework is also MLE



Methodology roadmap

Nature of Dependent Variable	Methodology
Continuous	OLS
Discrete 1/0	Probit/Logit
Multivalued Discrete, with a value order [i.e. lower numbers are 'better' or worse']	Ordered Probit/Logit
Categorical labels	Multinomial Probit/Logit